jupyter

Jupyter in HPC

Feb 28th, 2018

Matthias Bussonnier

bussonniermatthias@gmail.com
GitHub: @carreau
Twitter: @mbussonn

About Me

Matthias Bussonnier

- A Physicist/Bio-Physicist
- Core developer of IPython/Jupyter since 2012
  - Co-founder, and Steering Council member
- Post doctoral Scholar on Jupyter at BIDS

Demo(s)

Just to find bugs and make things crash

Webinar & Outline

- This webinar will be in 3 parts
  - Overview of what is Jupyter + HPC
  - Use case : Suhas Somnath
  - Use case : Shreyas Cholia
- Outline Part 1
  - From IPython to Jupyter
  - What is Jupyter
  - Jupyter Popularity
  - Some Jupyter Usage

From IPython to Jupyter

- 2001: Fernando Perez Wrote "**IPython**"
  - Create IPython for Interactive Python with prompt number, gnu plot integration
  - Replace a bunch on perl/make/C/C++ files with only Python.
- 2011: QtConsole
- 2012: Birth of current **Notebook** (6th prototype)
  - Make IPython "network enabled"
  - Made possible by mature web tech.
- 2013: First non-Python (**Julia)** kernel
- 2014: we **renamed** the Python-Agnostic part to **Jupyter**.
- 2018: several millions users & **JupyterLab** released

What is Jupyter

- Mainly Known for **The Notebook**
  - Web server, a web app, load .ipynb (json), containing code, narrative, math and results.
  - Attached to a **Kernel** doing computation.
- Results can be:
  - Static (Image)
  - Interactive (client-side scoll/pan/brush)
  - Dynamic (Call back into Kernel)

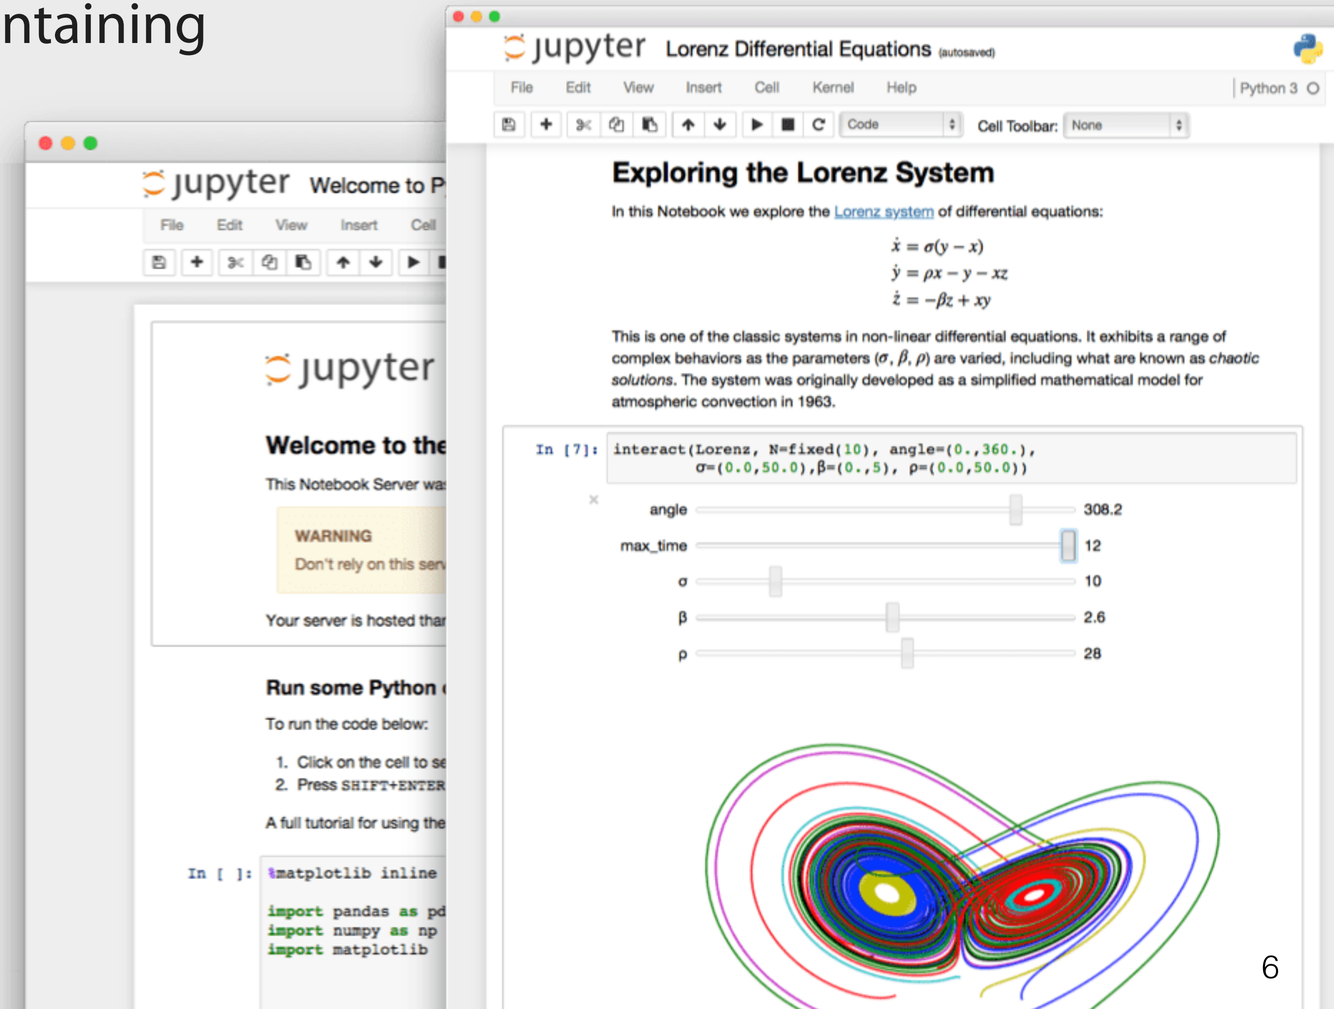

Focused on Exploratory Programming

- IPython was designed for exploratory programming, as a REPL (Read Eval Print Loop) and grew popular, especially among scientist who loved it to explore.

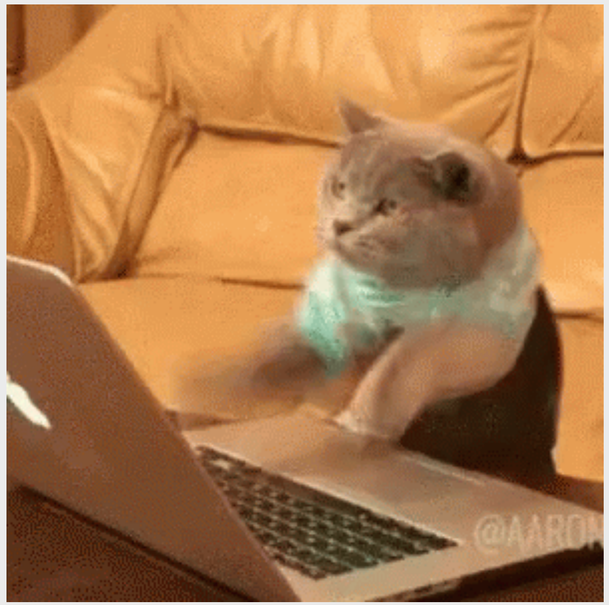

"IPython have weaponized the tab [completion] key"
– Fernando Pérez

Open Organisation

- Organisation with Open Governance (https://GitHub.com/jupyter/governance)
- Funded by Grants and Donations, and Collaborations

NUMFOCUS

OPEN CODE = BETTER SCIENCE

GORDON AND BETTY MOORE FOUNDATION

Alfred P. Sloan FOUNDATION

THE LEONA M. AND HARRY B. HELMSLEY CHARITABLE TRUST

Protocols and Formats

- Jupyter is also a set of **Protocols and Formats** that reduce the **N-frontends × M-backends** problem to a **M-Frontends + N-backends**,
  - Open, Free and Simple.
    - JSON (almost) everywhere
    - Notebook document format,
    - Wire protocol
  - Thought for Science and **Interactive** use case.
    - Results embedded in documents no "Copy past" mistake.
    - Scale from Education to HPC jobs.

Ecosystem

Frontends: Notebook, JupyterLab, CLI, *Vim, Emacs, Visual Studio Code, Atom, Nteract, Juno...*

Kernels: Python, *Julia, R, Haskell, Perl, Fortran, Ruby, Javascript, C/C++, Go, Scala, Elixir... 60+*

Building Blocks: Nbformat, JupyterHub, Kernel Gateway...

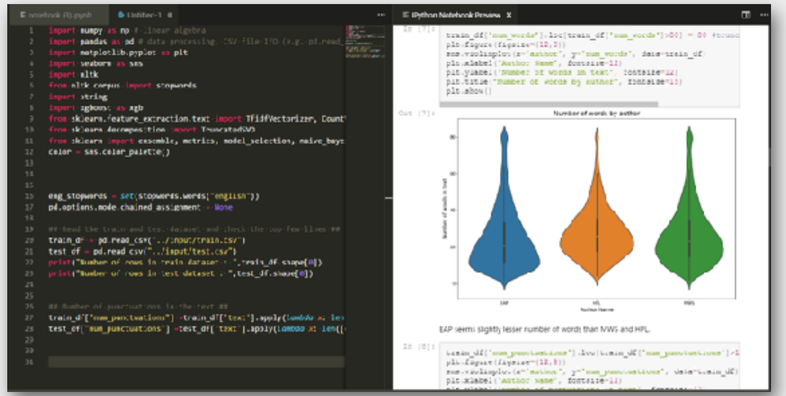

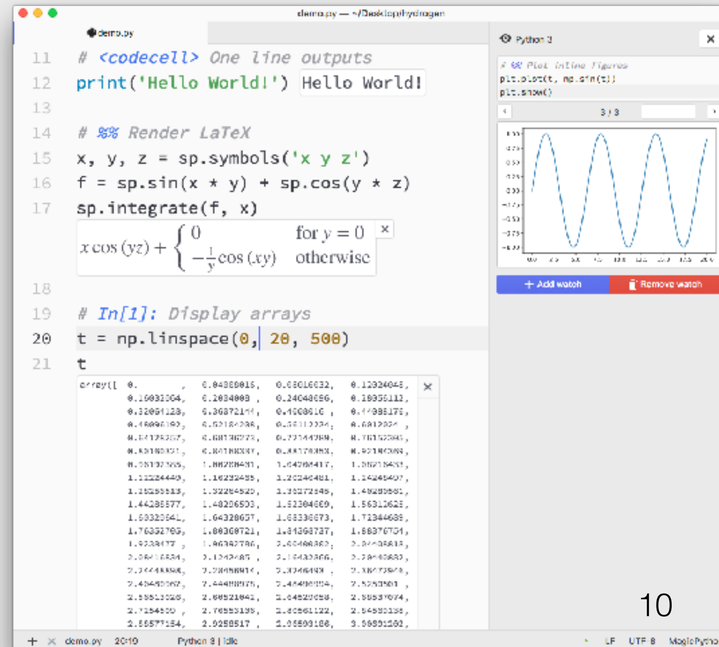

JupyterLab

- Extends the notebook interface

  with text editor, shell, ...etc

- is it and IDE ?

- If by I you mean Interactive,

  then yes

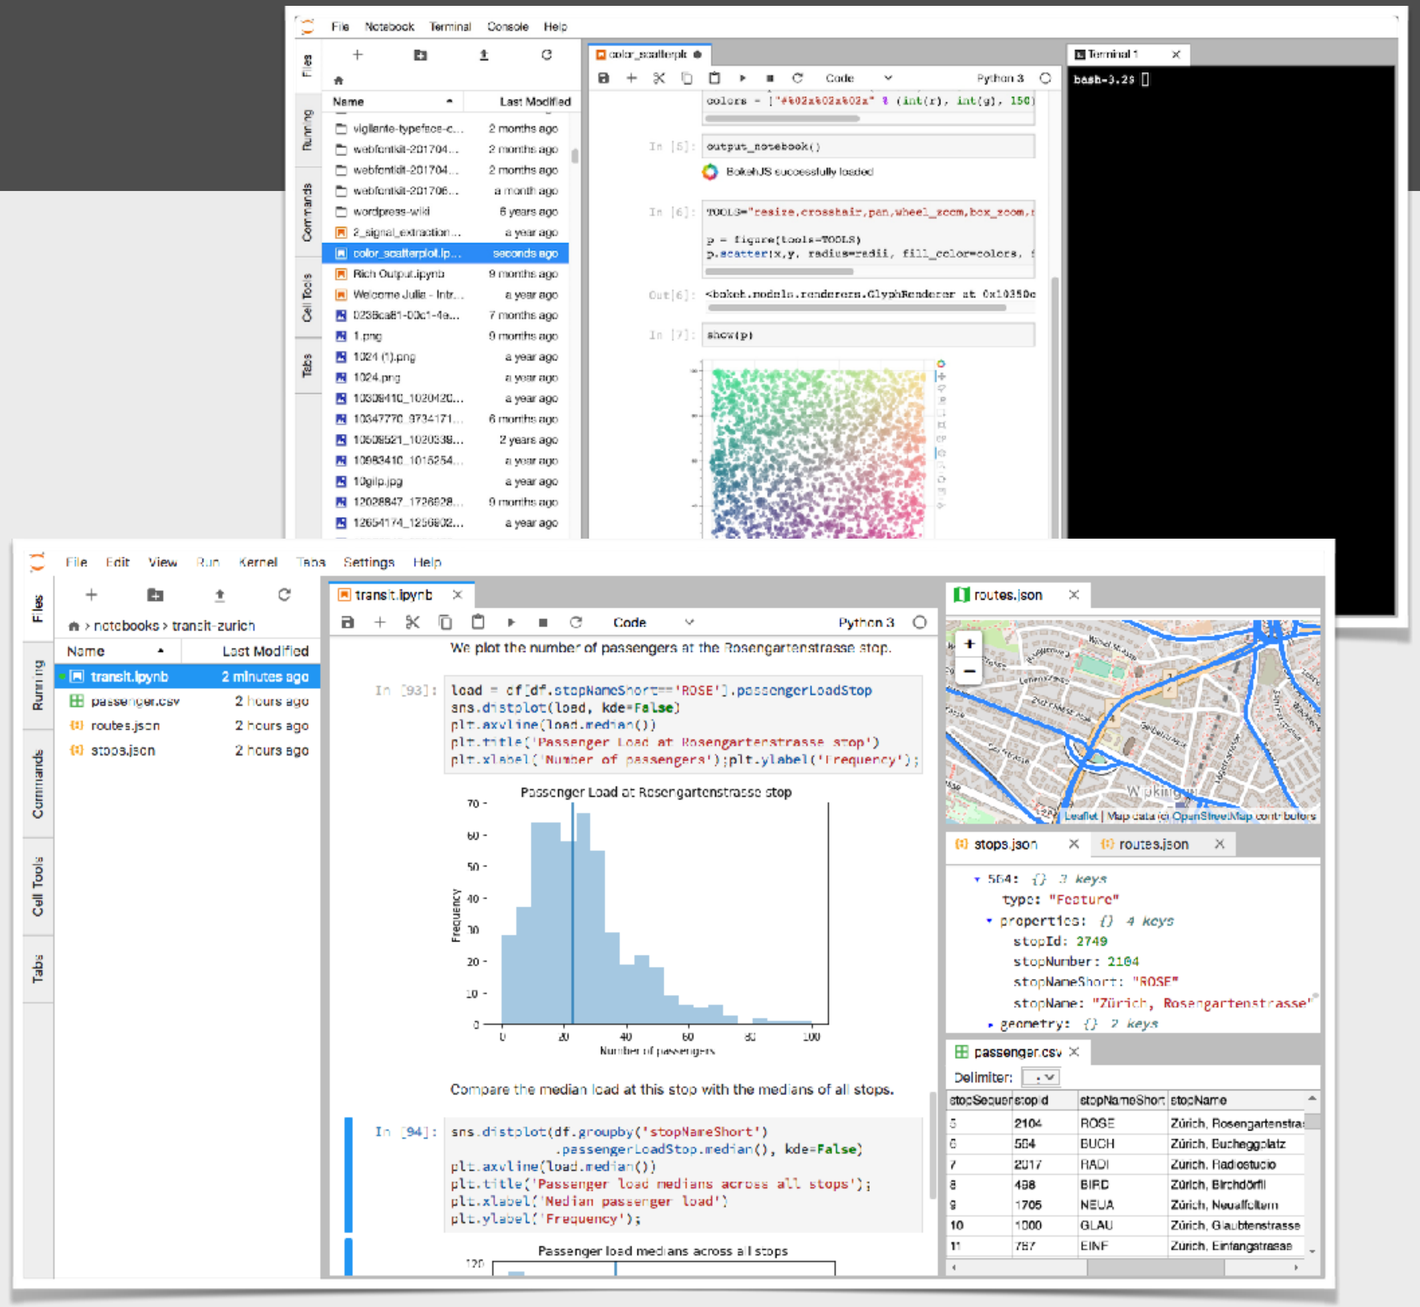

https://blog.jupyter.org/jupyterlab-is-ready-for-users-5a6f039b8906

Popularity

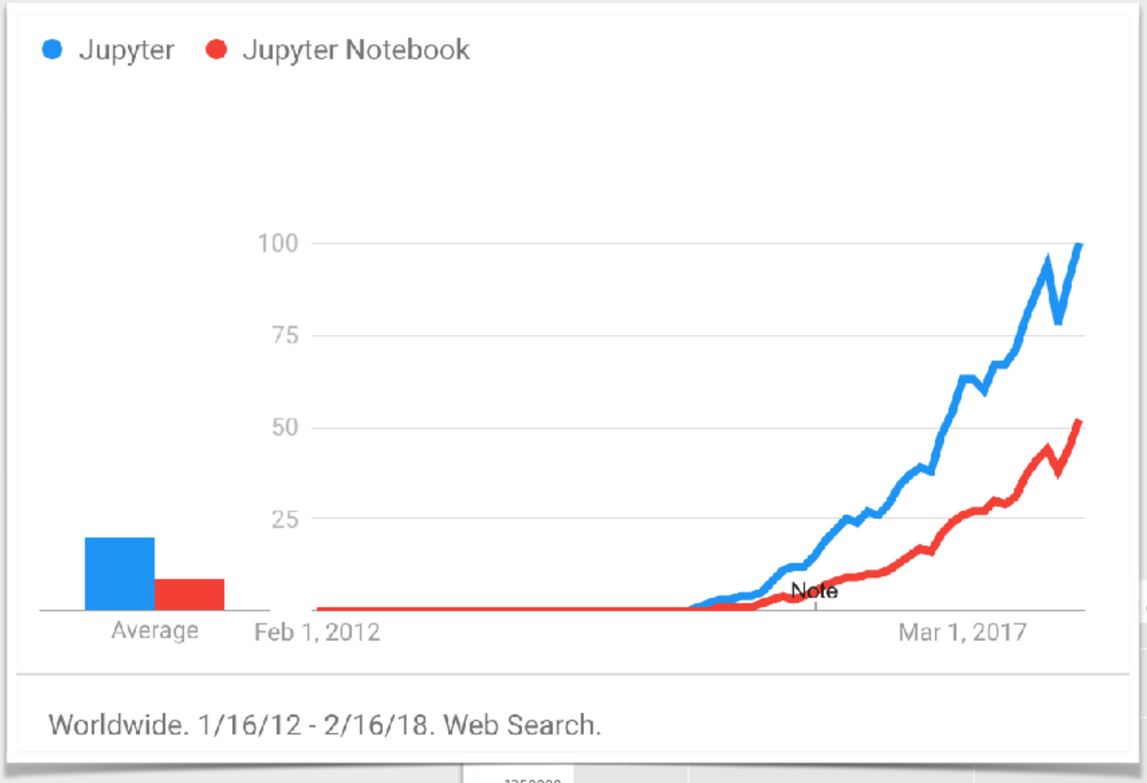

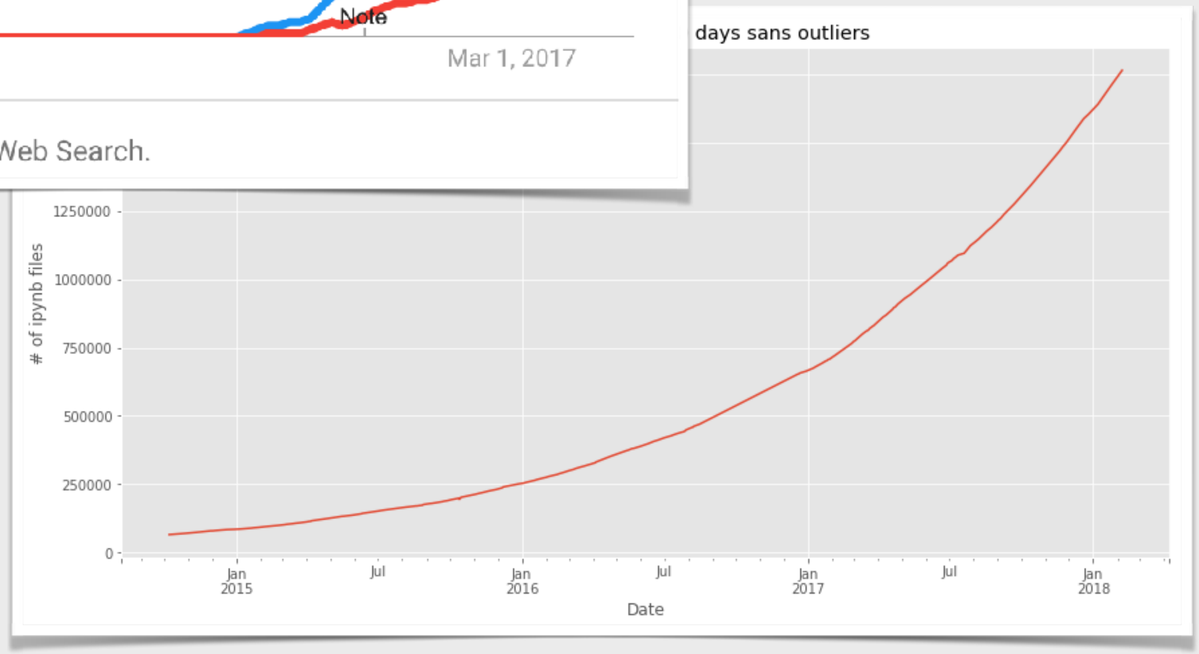

https://github.com/parente/nbestimate

Popularity

Interactivity

- Coding is not the end goal of most of our users. A simple, single tool, with friendly interface helps.
- Persisting kernel state allows to iterate only on part of an analysis.
- Notebook interface give the interactivity of the REPL with the edit-ability and linearity of a script with intermediate result.

  Aka "Literate Computing”

Popularity

Separation of states

- Computation, narrative/visualisation in different processes.
  - Robust to crashes
  - Can "Share" and analysis / notebook without having to “rerun"
  - Trustworthy (No copy-past issues).
- Cons:
  - Understanding that document/kernel can have different states can be challenging.
  - Notebook format is not as widespread as others.

Popularity

Network enabled / web based

- User love fancy colors and things moving. Using D3 and other

  - dynamic libraries are highly popular
  - Usable by novices and power-users
  - Users w/ different expertise (Numerical Methods, Visualization,...)
- Seamless transition to HPC: Kernel Menu > Restart on Cluster
- Document persist if code crash.
- Can be Zero-Installation (See JupyterHub).
  - A web browser is all you need.

JupyterHub

- Multi-users Jupyter deployment
  - Not (Yet) Realtime collaboration
- Each user can get their own process/version(s)/ configuration(s)
  - Hooks into any Auth
  - Only requires a browser
- Not limited to running Jupyter (e.g. work with RStudio, OpenRefine...)

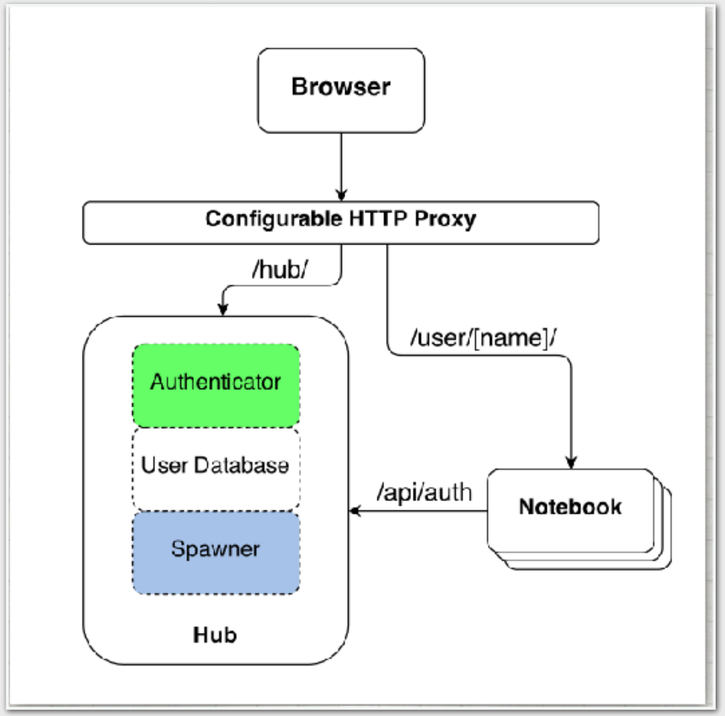

Use Cases

HPC

- Batch Jobs
  - You can run notebook “headless”
  - Parametrized notebook as “reports” you can interact with later
- Interactive Cluster.
  - Run a Hub (hook into LDAP/PAM...)
  - Run notebook servers on a Head node
  - Run Kernels on head Node/fast queue
  - Extra Workers (e.g. dask) on Batch queue/cluster.

HPC deployment

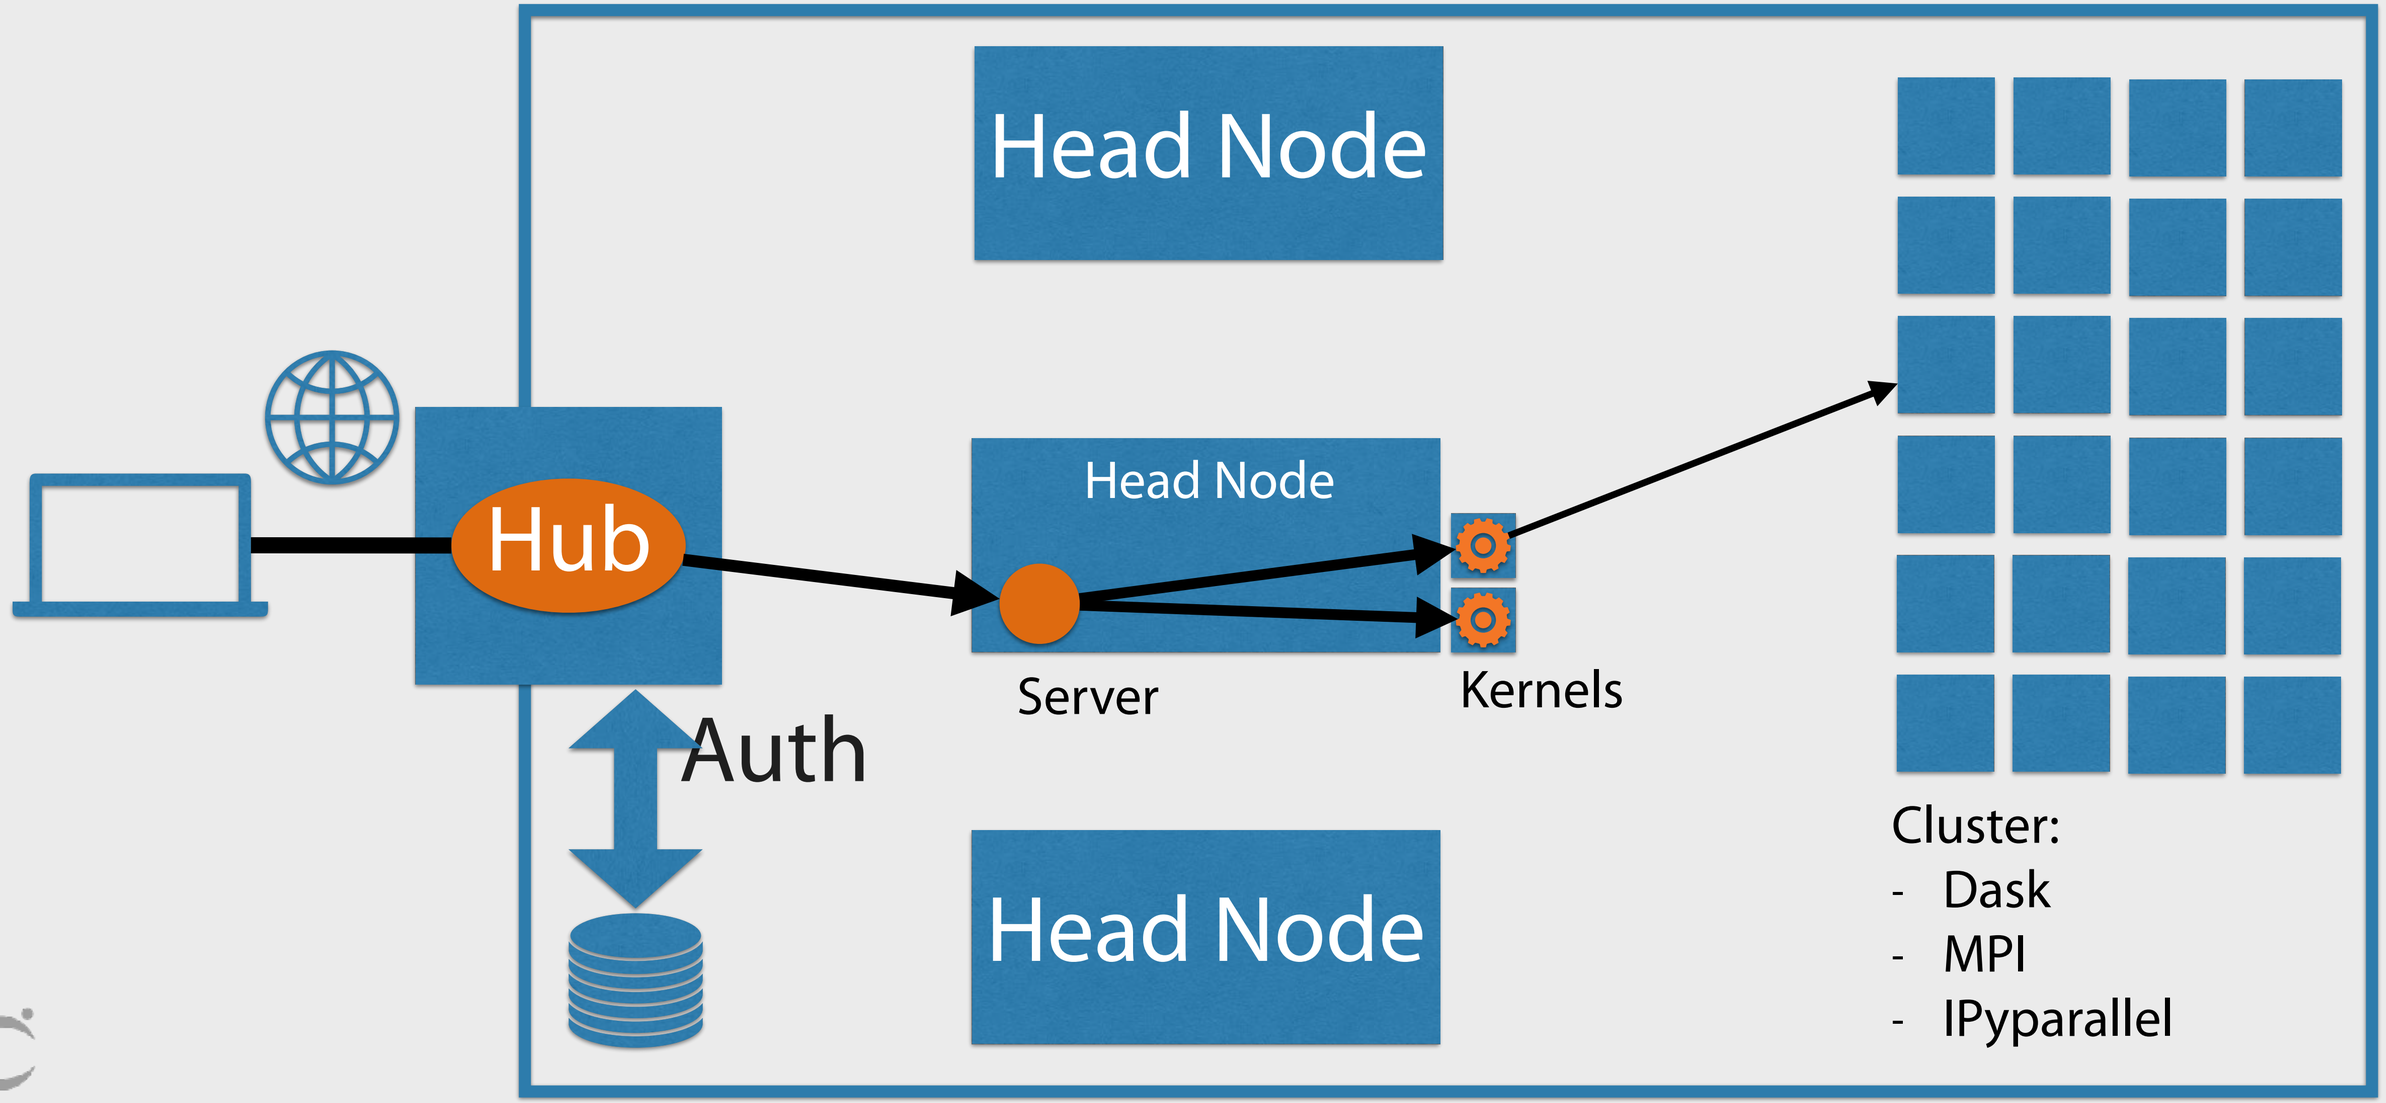

HPC Misconceptions

We need to run JupyterHub on the cluster: No

- Hub, Server, Kernels, (and Workers):
  - Do not have to be on the same machine
  - Do not have to use the same environment

A Kernel is a (single) language: No

- A Kernel is a preconfigured computation environment. It can be:
  - A queue, a hardware resource (GPU, SSD...), A location (like a beam-line)
- Example of Python, Cython, Julia, R, Fortran, Rust, C calling each other in same notebook

Every User have the same environment: No/No

- Kernels and notebook server can be configured independently
- Subset of users could use different server versions w/ different extensions.

JupyterHub is Limited to Jupyter: No

JupyterHub Can run RStudio, Open Refine.

Danger !

Despite Notebook being great, some limitations:

Most if not all **document state** is **in your browser** !

- Watch out for flaky network connections !
- Do Not close your Laptop Lid*/Tab* !

Workaround:

- Wrap computations (especially long), in Futures
- Use Caching.

Interrupting in compiled code is hard.

Large outputs/notebooks can crash the browser

*recent versions have buffering if client disconnected properly

Some Jupyter Usage

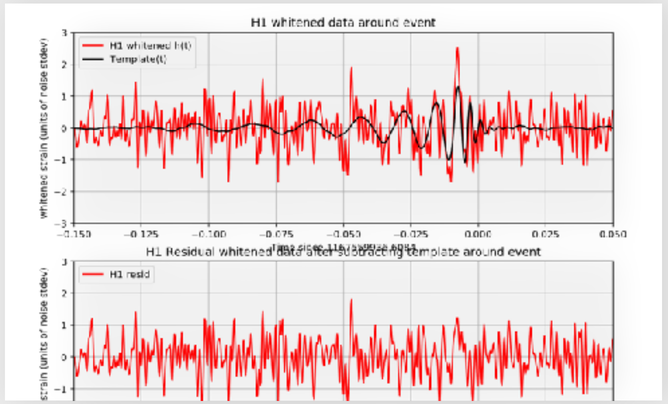

Ligo

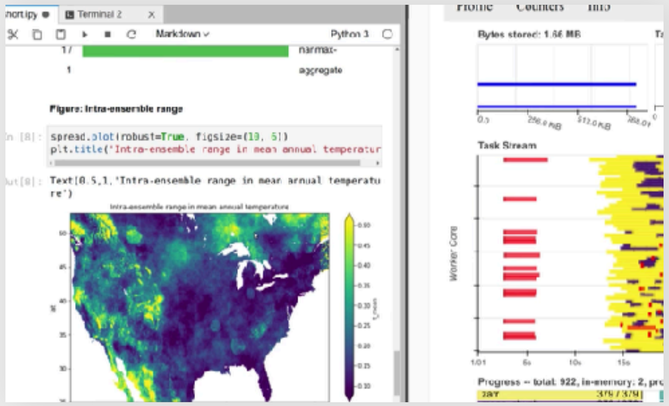

Pangeo

Cern’s SWAN

Ligo

- Some events analysis with Jupyter
- Subset of data + env put online
- Run the analysis yourself on Binder[1] and listen to the waves

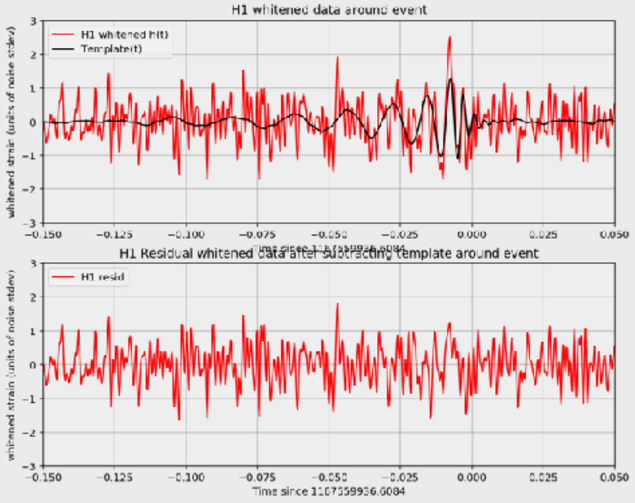

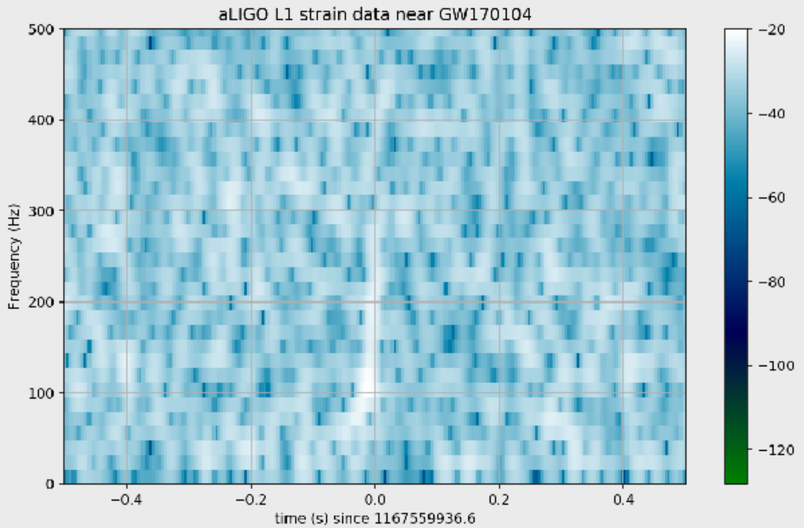

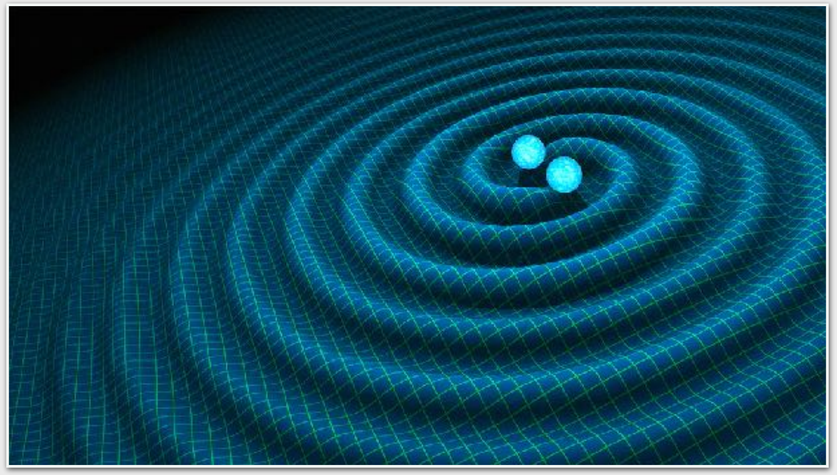

[1] https://github.com/minrk/ligo-binder

Pangeo (pangeo-data.github.io)

- Effort from Atmosphere / Ocean / Land / Climate (AOC) science community
- unified effort
- Cloud based
- Recent Technologies
  - Dask, Jupyter

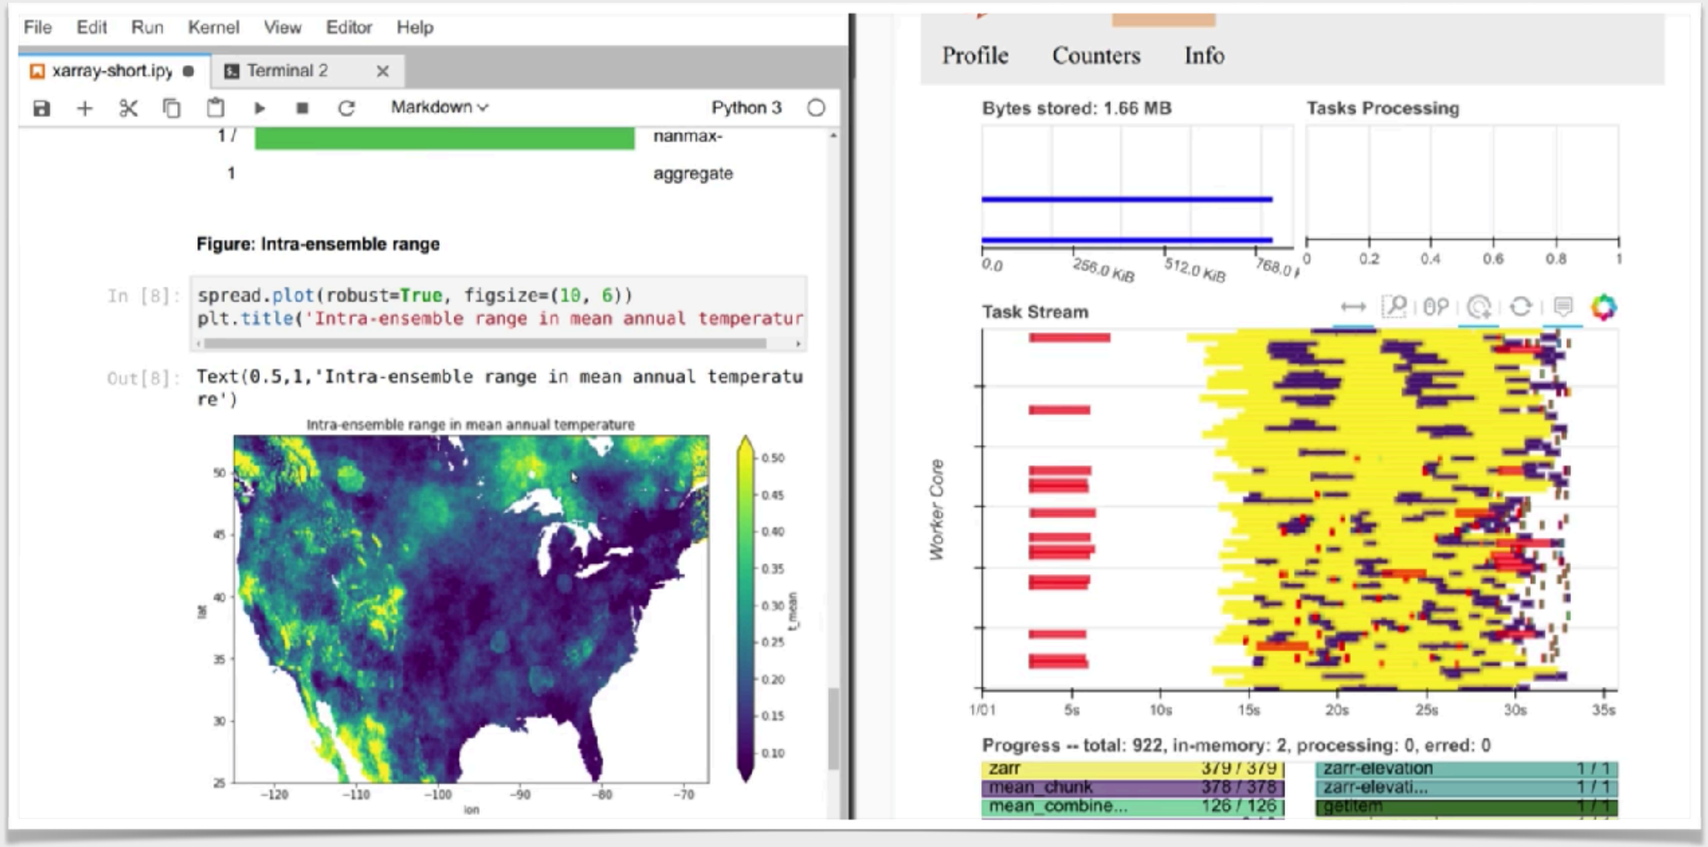

Matt Rocklin Blog post on pangeo-data.github.io

Cern Swan (swan.web.cern.ch)

- Share platformed for Data Analysis
- Sync W/ $HOME directory
- 0-install
- Share Data
- Provide example gallery with 1-click-fork

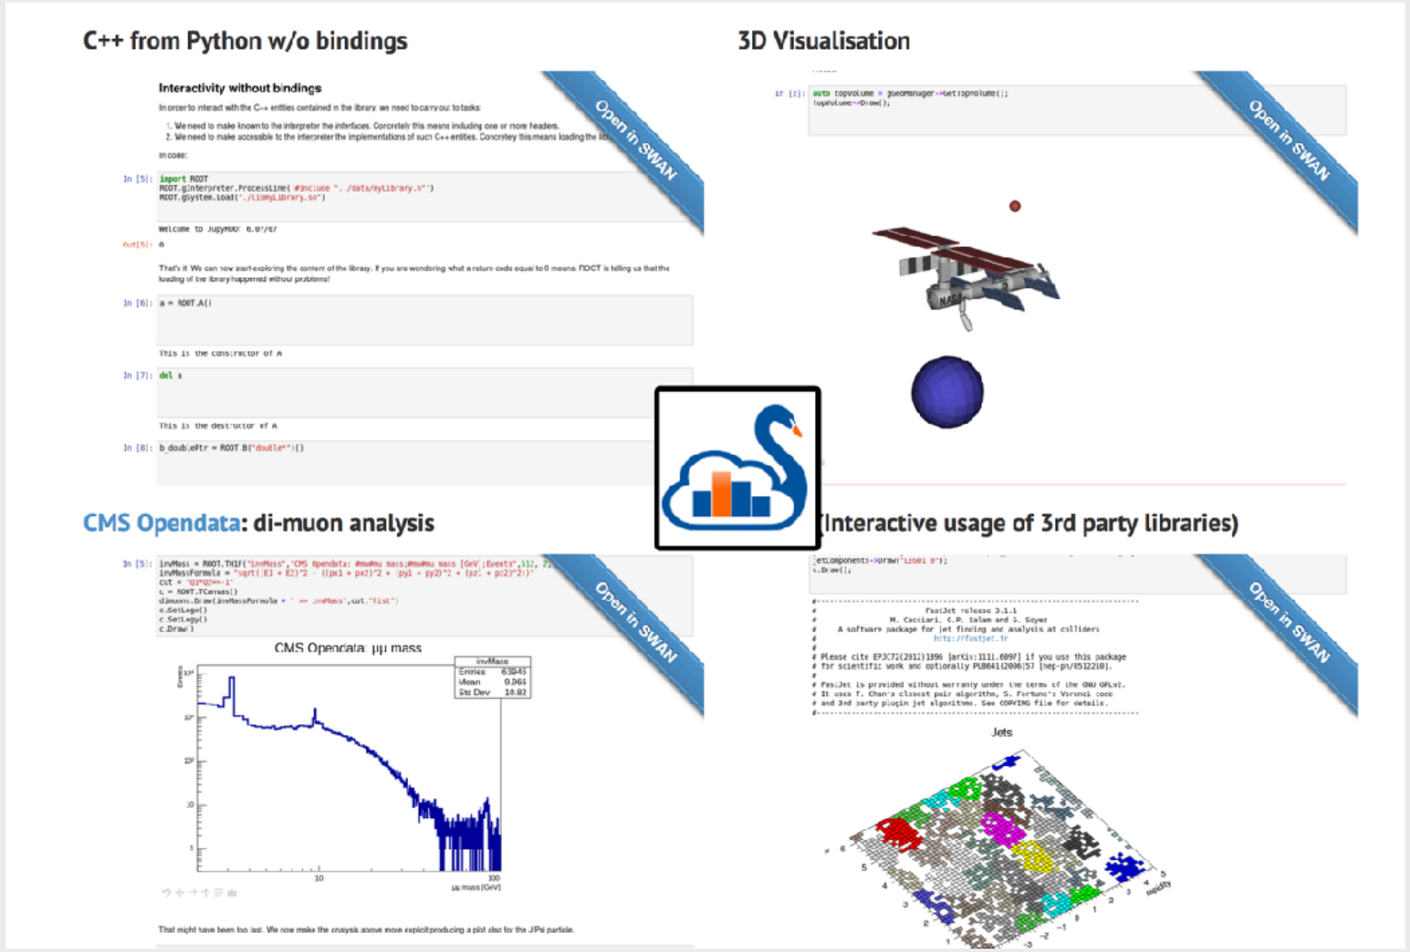

The Shape of Things to come

The Shape of Things to come

Classic Notebook -> JupyterLab transition

- Stabilisation
- Transfer of extensions
- Collaboration:
  - Google retired Real-time API
  - Who “executes” problem
- Long Running Jobs

JupyterHub

- Horizontal (and Vertical) Scaling
- Audits APIs (Hippa Compliance ?)
- “Federation” (binder) / Intercommunication

CFP- Ends March 6th

Question(s)
while we change
speakers ?